

PROFESSIONAL SUMMARY:

Data Engineer with 5+ years in analytics and pipeline development across retail, e-commerce, and pharma. Built scalable ELT pipelines (Python, PySpark, dbt) on AWS and Snowflake, delivering data models that drive revenue, improve accuracy, and reduce defects. Proven track record of recovering \$1M+ GMV and cutting manual reporting by 90%.

SKILLS

Programming Skills: SQL, Python, PySpark, Spark SQL, Scala, Airflow, dbt, Terraform, AWS (Glue, Lambda, S3, Redshift), JavaScript, Docker, Git, REST APIs

Theoretical Knowledge: ETL | ELT | Snowflake | dbt | Redshift | Athena | Data Modeling | Star Schema | CDC | Power BI | Tableau | QuickSight | A/B Testing | BigQuery | Kafka | Data Quality | ML | DAX | Window Functions

Education: Texas State University **Majors:** Data Analytics and Information Systems | GPA: 3.7 **Graduation Date:** 2024

EXPERIENCE

Real Value Products – Value RX

Data Engineer

*San Antonio, TX
June 2024- Present*

- Built end-to-end ELT pipelines in Python & PySpark to ingest vendor APIs and web-scraped catalog data into AWS S3; implemented Bronze/Silver/Gold medallion architecture with Spark SQL transformations – improving data accuracy by ~20%.
- Modeled retail analytics data warehouse using dbt on Snowflake with star/snowflake schemas (FactSales, FactInventory; DimProduct, DimVendor) – implemented incremental materializations and dbt tests to ensure data quality at scale.
- Optimized SQL Server stored procedures and batch loads (MERGE, window functions, CDC) with Snowflake's zero-copy cloning for dev/test environments – reduced heavy report queries from minutes to <10 seconds.
- Implemented automated data-quality checks using dbt tests and Python reconciliation jobs (nulls, type mismatches, UPC padding) – reduced listing defects by ~30% through exception logging and alerting.
- Designed Power BI executive dashboards with RLS, custom DAX measures (GM%, ASP, Buy-Box rate, OOS%, return rate) enabling weekly revenue/margin forecasting and SKU-level actions across eBay, Walmart, Amazon, Shopify.

Texas State University

Data Engineering Assistant

*San Marcos, TX
June 2022- May 2024*

- Built reusable SQL/Python pipelines to extract and transform data from Banner/Canvas and campus surveys – cut weekly report prep time by ~60% through automated ETL processes.
- Developed interactive Tableau/Power BI dashboards for enrollment, retention, and course performance – enabled deans/advisors to track KPIs in real time, reducing conflicting metrics across departments.
- Conducted statistical analyses (A/B tests, regression) to evaluate tutoring programs; surfaced drivers of GPA lift and course pass rates leading to data-driven program improvements.
- Implemented data-quality checks (null/duplicate audits, validation rules) and created lightweight data dictionary to standardize metrics (FTIC, DFW rates) – improved report accuracy and audit readiness.

Amazon

Business Intelligence Analyst (Data Engineering)

*Bengaluru, KA
March 2020 – August 2022*

- Built automated vendor health dashboard (Athena + QuickSight) aggregating case logs, chargebacks, and defect codes from S3 – reduced manual reporting time by 90% and cut vendor SLA breaches by 28%.
- Created Python ETL pipelines (Lambda + Glue) normalizing SP-API feeds (orders, inventory, pricing) across 200+ vendors – improved data freshness from daily to hourly, increased anomaly detection precision by 35%.
- Designed SQL models in Redshift to track Buy Box inhibitors (pricing gaps, content defects, OOS risk) – partnered with Category Managers to drive interventions lifting Buy Box win rate by 7.4% and recovering ~\$1.2M GMV.
- Led root-cause analyses on PO-to-ship defects using cohort and Pareto analysis; implemented vendor playbooks

reducing chargeback incidence by 22% and improving on-time-ship by 11 p.p. in 2 quarters.

- Stood up alerting pipeline (CloudWatch + SNS) for price parity and MAP violations with dynamic thresholds – cut time-to-detection from 12h to <30m, preventing ~\$400k in margin leakage annually.

PROJECTS

- Retail Analytics Pipeline – dbt + Snowflake
Built end-to-end analytics pipeline using dbt Core on Snowflake; implemented incremental models, schema tests (unique, not_null, relationships), and auto-generated documentation. Deployed via dbt Cloud with CI/CD integration for automated testing on PRs.
- E-Commerce Data Warehouse – Snowflake + AWS
Designed dimensional data warehouse (star schema) on Snowflake with staging, intermediate, and mart layers. Ingested data from S3 using Snow pipe for real-time streaming; implemented role-based access control and query optimization with clustering keys.

CERTIFICATIONS

- AWS Certified Solutions Architect | Google Data Analytics Professional Certificate | NoSQL Essential Training